

Understanding exchange network dynamics with Python

Omer Yuksel

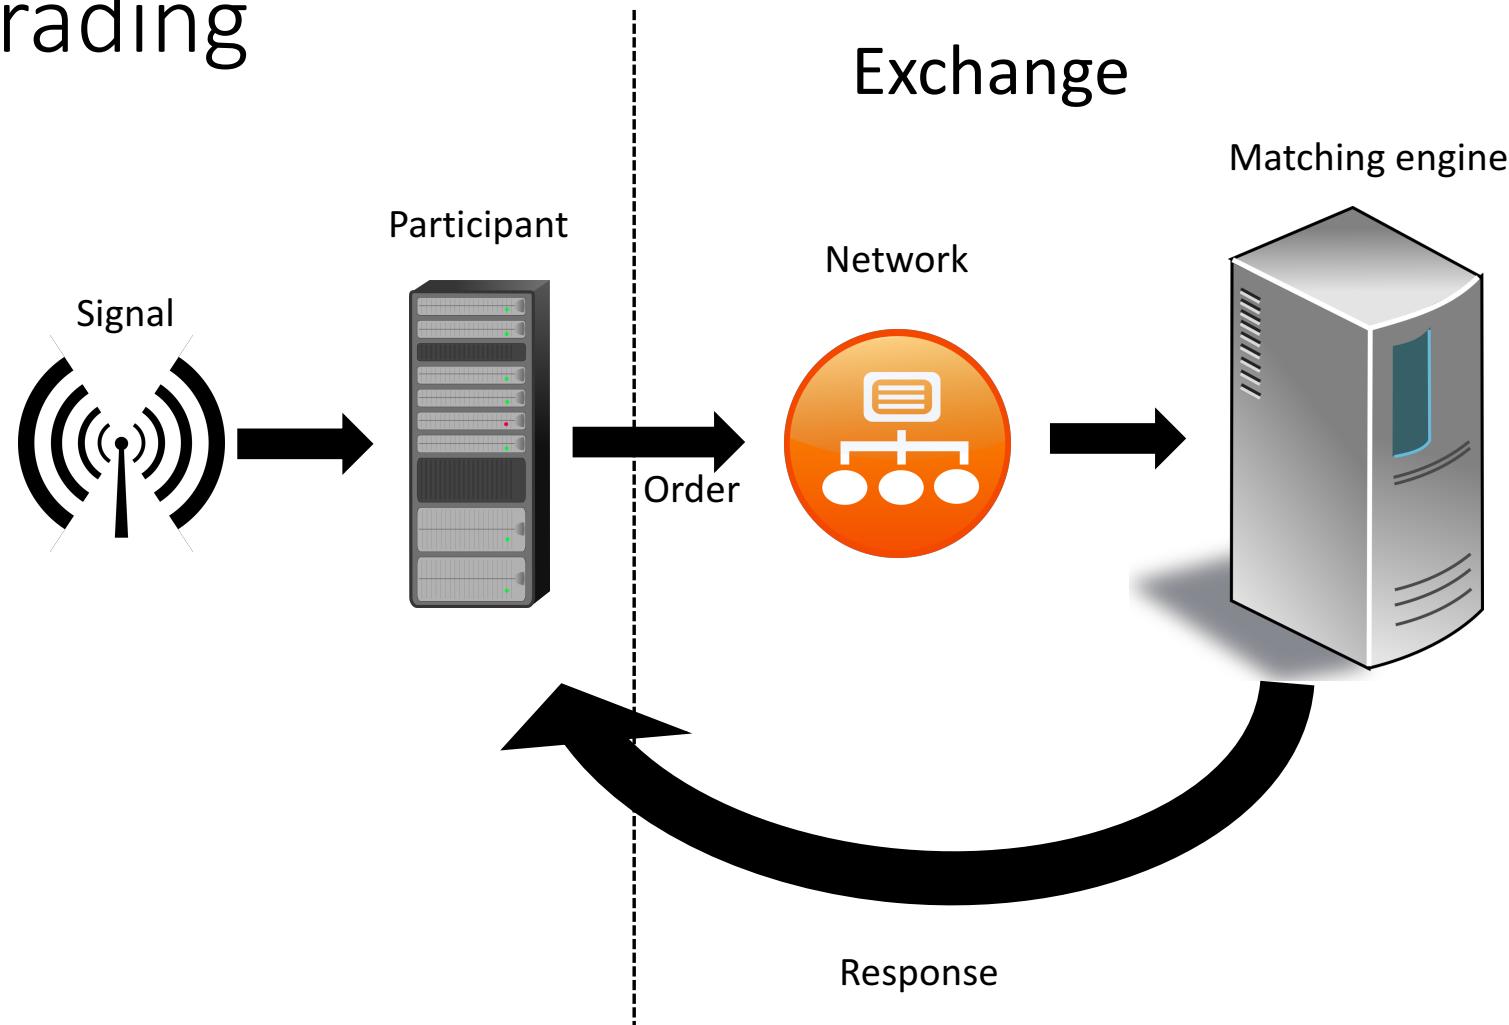
Introduction

- Performance Analyst at IMC
- Technology-driven trading
- Amsterdam, Chicago, Sydney



www.imc.com

Trading

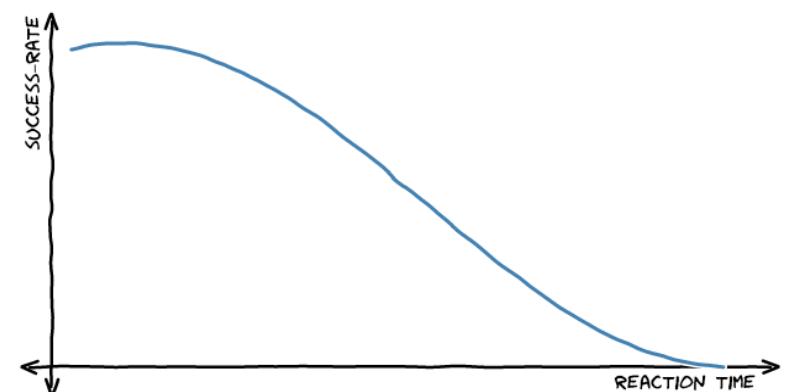


Exchange

Matching engine

Questions

- How often are we queuing?
- How much does queuing matter?
- How much should we improve our reaction time?

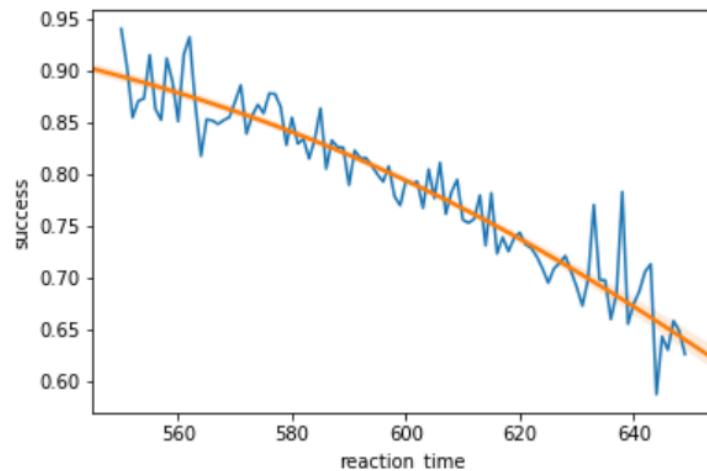


Methodology

Combine data with knowledge

- Data
 - Success/Failure
 - Reaction time
 - Time to reach matching engine
- Knowledge
 - Network structure
 - Components' behavior
 - Queuing effects

	ts	reaction_time	time_to_me	success	total_time_to_me
0	2019-06-03 09:00:00.042956193+02:00	606	5085	False	5691
1	2019-06-03 09:00:00.207588123+02:00	647	4500	False	5147
2	2019-06-03 09:00:00.276477343+02:00	633	4469	True	5102
3	2019-06-03 09:00:00.306361141+02:00	577	4459	True	5036
4	2019-06-03 09:00:00.389016054+02:00	599	4349	True	4948
5	2019-06-03 09:00:00.426910915+02:00	578	4515	True	5093
6	2019-06-03 09:00:00.552713100+02:00	572	4472	True	5044
7	2019-06-03 09:00:00.783978162+02:00	566	4680	True	5246
8	2019-06-03 09:00:00.870361884+02:00	606	4457	True	5063
9	2019-06-03 09:00:00.948479263+02:00	563	4394	True	4957

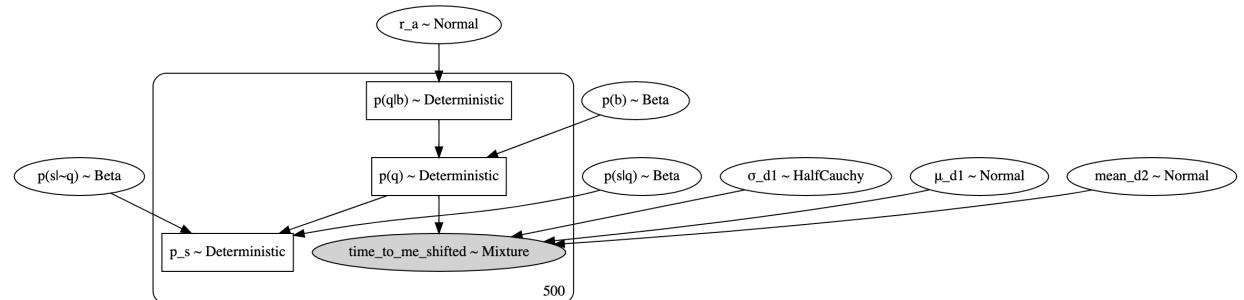


Approaches

- Probabilistic programming with PyMC3
 - Distributions and parameters
- Discrete-event simulations with SimPy
 - Code

PyMC3

- Bayesian probabilistic programming module
- Variables
 - Random
 - Free
 - Observed
 - Deterministic



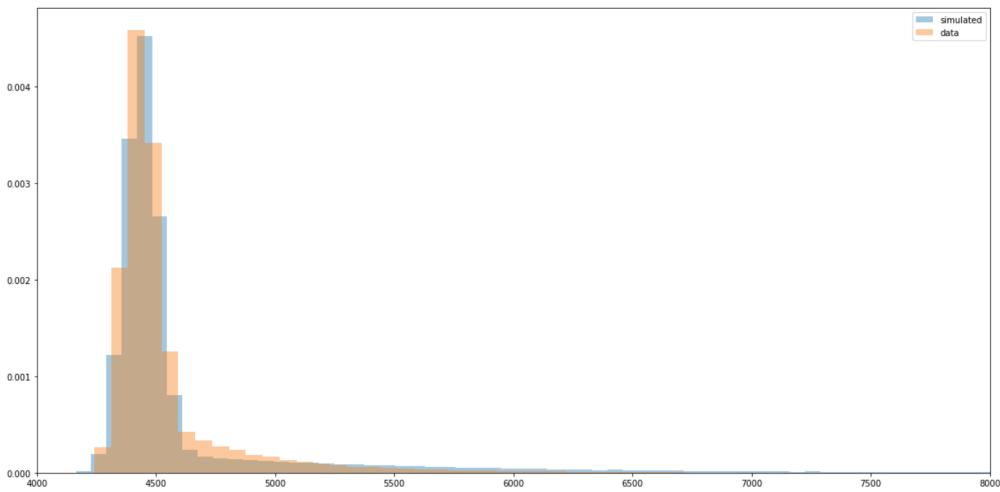
High-level model

- Hyperparameters
 - Performance: $perf \sim Beta$
 - Probability of bursts: $p_b \sim Beta$
 - Probability of success despite queuing: $p_{s_a} \sim Beta$
- Model parameters
 - Prob. queuing with/without bursts
 - Prob. success
 - Time to matching engine with/without queuing
- Observed
 - Time to matching engine
 - Reaction time
 - Success failure

Success: $s \sim Binomial(p_s)$
Reaction time: $r \sim Normal(\mu_r, \sigma_r)$
Time to matching engine: $d_{me} \sim Mixture(d_1, d_2, w)$
- Method: ADVI

Posterior sample

Time to matching engine:



Success rate:

Simulation = 78.8%

Data = 78.9%

Posterior sample

		mean	sd	mc_error	hpd_2.5	hpd_97.5
Reaction time	r_mu	599.846909	1.566378	0.016985	596.739331	602.869329
	r_sigma	22.180482	0.772145	0.008591	20.637166	23.673577
Queuing	p_q	0.191285	0.029576	0.000252	0.135996	0.251843
	p_b	0.255693	0.028668	0.000265	0.201732	0.313731
Success rate	p_s	0.788471	0.024746	0.000210	0.738779	0.835984
	p_s_a	0.424210	0.039147	0.000366	0.349996	0.502416

PyMC3 – What's next?

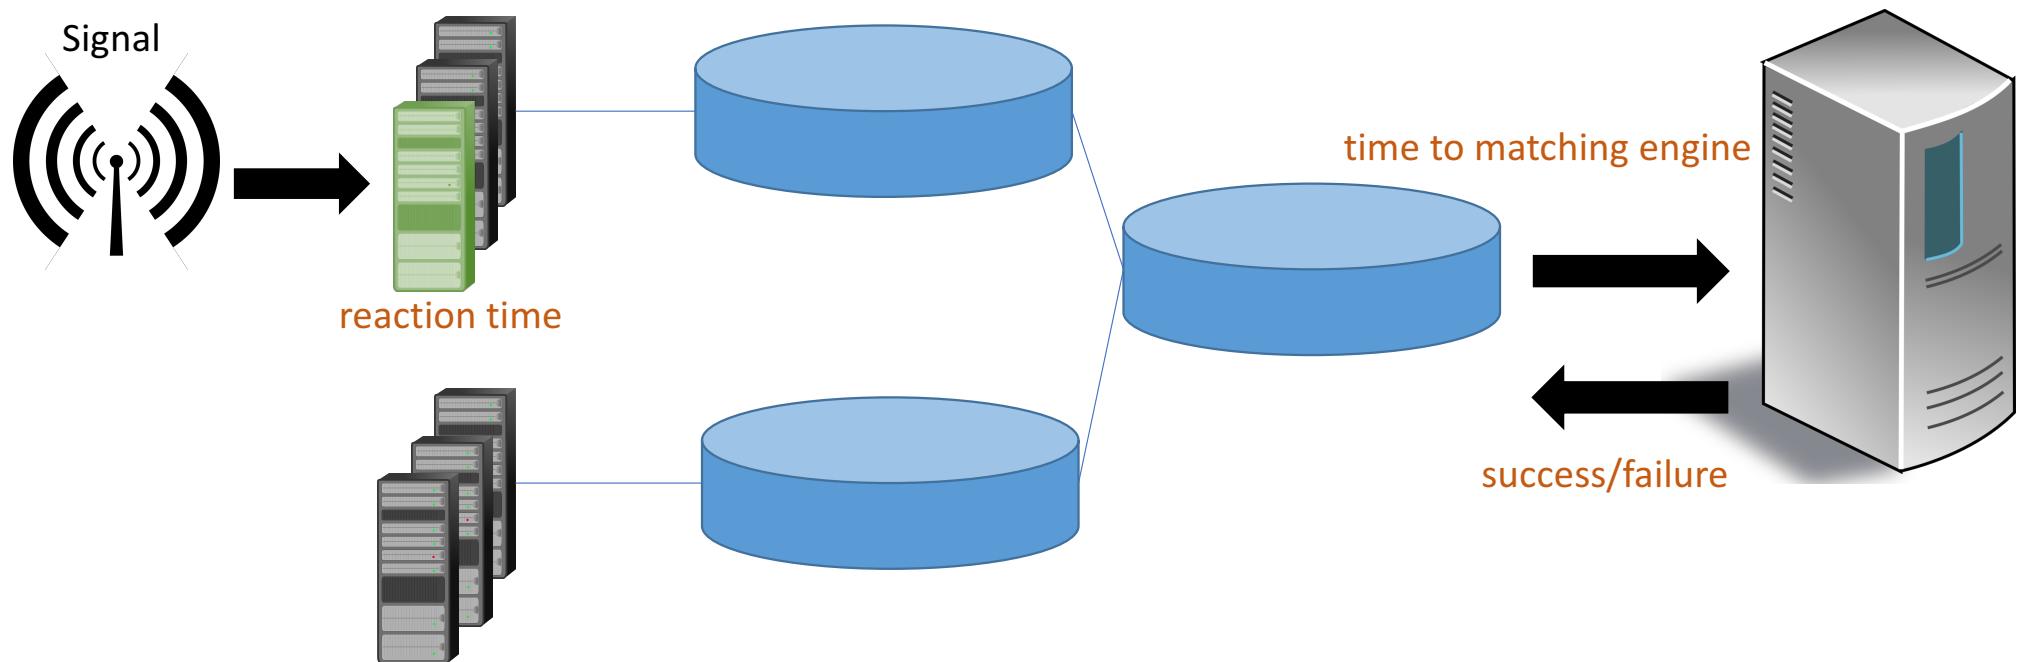
- Different models
 - Additive model for latencies
 - Explicit model for reaction window
- MCMC vs VI

SimPy – Discrete Event Simulation

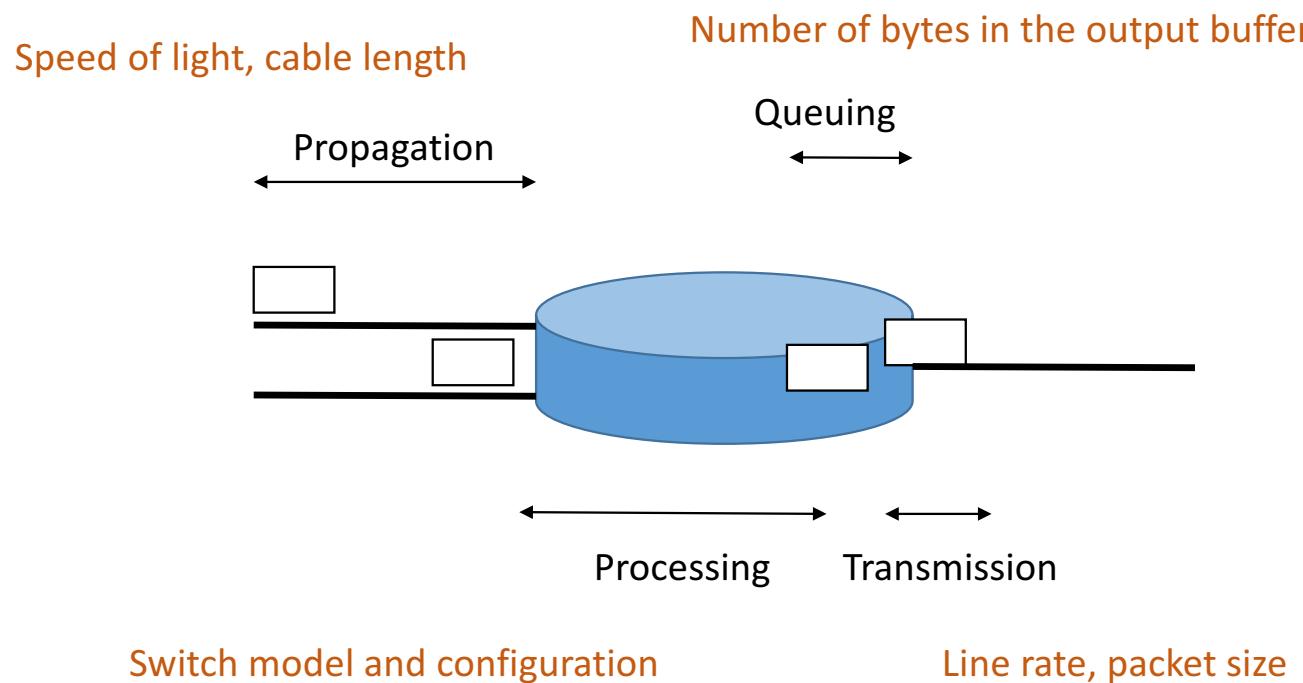
- Environment
- Event – “signal received”
- Process – participant, switch,
- Resources – queue, semaphore, ...

<https://simpy.readthedocs.io/en/latest/contents.html>

Network



Delays



Bursts and Queuing

- Bursts - multiple packets
- Queuing:
 - Transmission delays of the packets in front
 - Non-deterministic

Building blocks - Simulation

```
def run_simulation(duration):
    env = simpy.Environment()
    ...

    # Nodes
    matching_engine = MatchingEngine('matching_engine', env)
    switch = FIFOswitch('Switch1', env, matching_engine, SWITCH_TO_MATCHING_ENGINE)
    alice = SimpleParticipant('Alice', env, switch, TIME_TO_SWITCH)
    bob = SimpleParticipant('Bob', env, switch, TIME_TO_SWITCH)

    env.run(duration)
```

Building blocks - Switch

```
class FIFOswitch(NetworkNode):
    def __init__(self, name, env, destination, time_to_destination):=

    def run(self):
        while True:
            packet = yield self.incoming_packets.get()
            self.debug(f'{packet} arrived')
            self.env.process(self.process_packet(packet))

    def process_packet(self, packet):
        yield self.env.timeout(SWITCH_PROCESSING_DELAY)
        self.outgoing_packets.put(packet)

    def transmit_outgoing_packets(self):
        while True:
            packet = yield self.outgoing_packets.get()
            self.debug(=)
            self.send_to_destination(packet)
            yield (self.env.timeout(TRANSMISSION_DELAY))
```

Simulation output

```
DEBUG [100000000] Alice - Signal received.  
DEBUG [100000000] Bob - Signal received.  
DEBUG [100000500] Alice - Sent packet to switch (Switch1)  
DEBUG [100000500] Bob - Sent packet to switch (Switch1)  
DEBUG [100001500] Switch1 - Packet (Alice) arrived  
DEBUG [100001500] Switch1 - Packet (Bob) arrived  
DEBUG [100002470] Switch1 - Transmitting packet. Remaining #packets in queue: 1  
DEBUG [100002650] Switch1 - Transmitting packet. Remaining #packets in queue: 0  
DEBUG [100003970] matching_engine - Received Packet (Alice).  
DEBUG [100004150] matching_engine - Received Packet (Bob).
```

Expanding the model

- Non-deterministic switch
- Signal generator
- Bursts

Exchange model

```
env = simpy.Environment()

# Nodes
me = MatchingEngine('matching_engine', env)

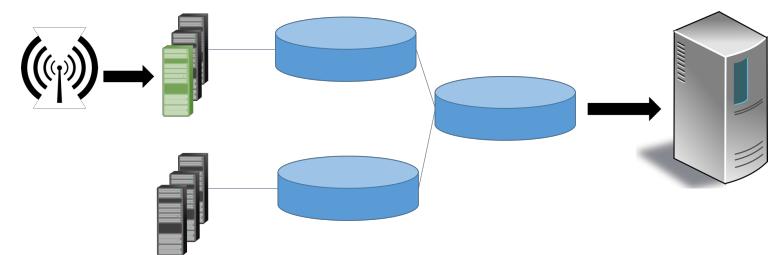
switch_2 = AlmostFIFOswitch('switch_2', env, me, 1000, switch_delay)

switch_11 = AlmostFIFOswitch('switch_11', env, switch_2, 500, switch_delay)
switch_12 = AlmostFIFOswitch('switch_12', env, switch_2, 500, switch_delay)

signal = BaseSignalGenerator('base_signal', env, signal_dist)

participant = Participant('participant', env, switch_11, 1000,
                         signal, reaction_time_dist_imc, window_dist)

burst = Burst('burst_1', env, switch_11, 1000, signal, reaction_time_dist_burst, num_reactions_dist)
burst_2 = Burst('burst_2', env, switch_12, 1000, signal, reaction_time_dist_burst, num_reactions_dist)
```



Exchange model

```
def run_simulation(duration, report_interval,
                    reaction_time_mean,
                    num_reactions_mean,
                    burst_reaction_offset,
                    window_size_loc):
    ...

    # Known distributions
    signal_dist = st.expon(scale=10 ** 8)
    switch_delay = st.truncnorm(loc=970 - REORDER_WINDOW, scale=50, a=-2, b=2)

    # Parameterized distributions
    reaction_time_dist_imc = st.truncnorm(loc=reaction_time_mean, scale=25, a=-2, b=2)
    num_reactions_dist = st.poisson(mu=num_reactions_mean)
    reaction_time_dist_burst = st.truncnorm(loc=reaction_time_mean + burst_reaction_offset,
                                              scale=70, a=-4, b=4)

    window_dist = st.lognorm(loc=window_size_loc, scale=10, s=7)

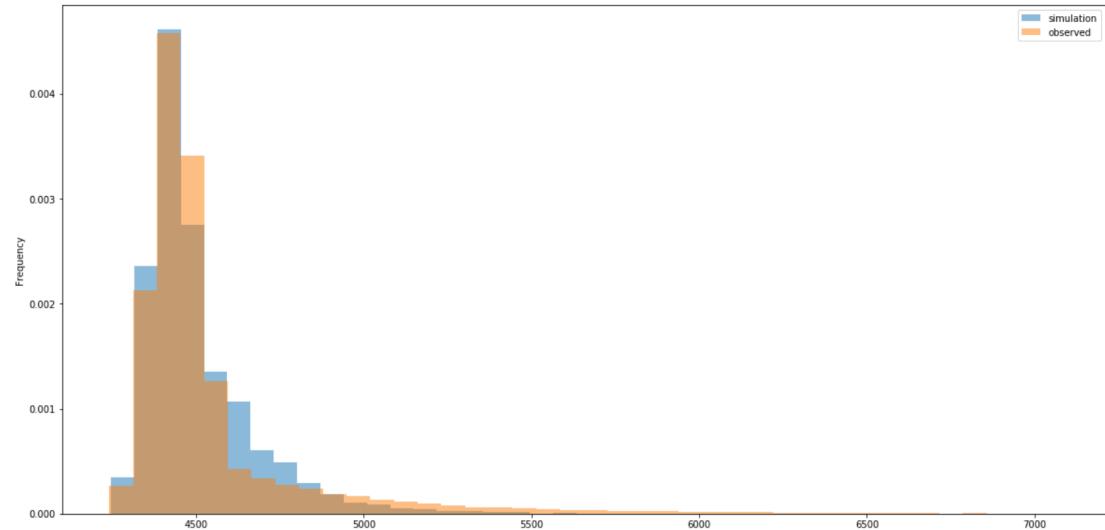
    ...
```

Fit model

- Parameters:
 - Burst behavior
 - #packets on average
 - Time offset
 - Reaction window
 - Average reaction time
- Cost function: Energy distance

Results

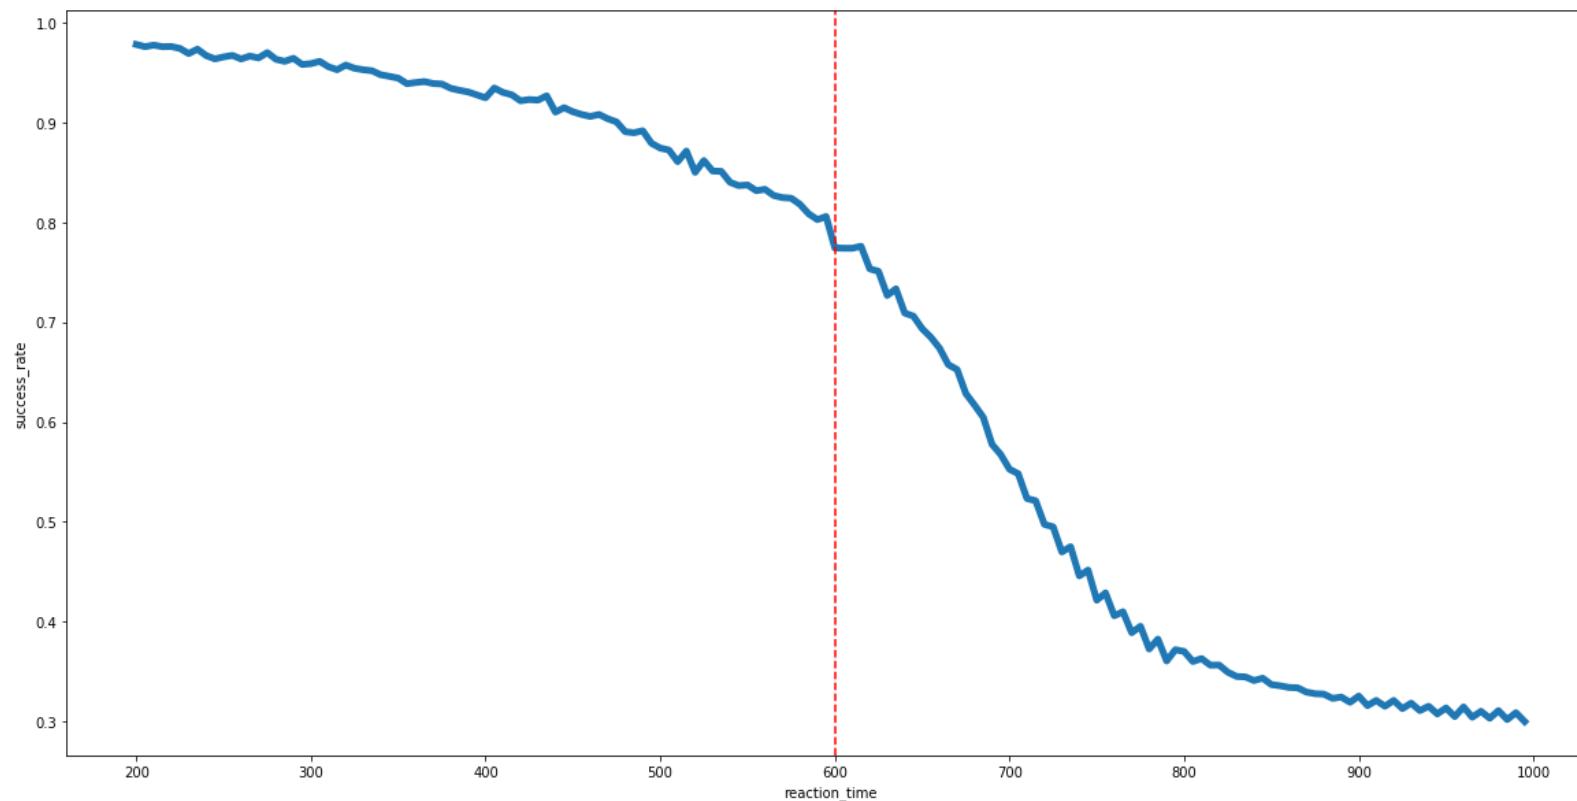
Time-to-matching engine:



Success rate:

- **Simulation: 78.4%**
- **Observed data: 78.9%**

Average reaction time vs success



When to use which?

- PyMC3
 - Concise
 - Many parameters
 - Fast with variational inference
- SimPy
 - Details
 - Discrete variables
 - Many moving parts
 - Scaling

Summary

- Incorporating knowledge
- Approaches
 - Discrete-event simulations
 - Probabilistic models

Questions



<https://github.com/omersyuksel/pydata-london-2019>